

# Real-time Personalization using Embeddings for Search Ranking at Airbnb

Mihajlo Grbovic

Airbnb, Inc.

San Francisco, California, USA

mihajlo.grbovic@airbnb.com

Haibin Cheng

Airbnb, Inc.

San Francisco, California, USA

haibin.cheng@airbnb.com

## ABSTRACT

Search Ranking and Recommendations are fundamental problems of crucial interest to major Internet companies, including web search engines, content publishing websites and marketplaces. However, despite sharing some common characteristics a one-size-fits-all solution does not exist in this space. Given a large difference in content that needs to be ranked, personalized and recommended, each marketplace has a somewhat unique challenge. Correspondingly, at Airbnb, a short-term rental marketplace, search and recommendation problems are quite unique, being a two-sided marketplace in which one needs to optimize for host and guest preferences, in a world where a user rarely consumes the same item twice and one listing can accept only one guest for a certain set of dates. In this paper we describe Listing and User Embedding techniques we developed and deployed for purposes of Real-time Personalization in Search Ranking and Similar Listing Recommendations, two channels that drive 99% of conversions. The embedding models were specifically tailored for Airbnb marketplace, and are able to capture guest's short-term and long-term interests, delivering effective home listing recommendations. We conducted rigorous offline testing of the embedding models, followed by successful online tests before fully deploying them into production.

## CCS CONCEPTS

• **Information systems** → Content ranking; Web log analysis; Personalization; Query representation; Document representation;

## KEYWORDS

Search Ranking; User Modeling; Personalization

## ACM Reference format:

Mihajlo Grbovic and Haibin Cheng. 2018. Real-time Personalization using Embeddings for Search Ranking at Airbnb. In *Proceedings of The 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining, London, United Kingdom, August 19-23, 2018 (KDD '18)*, 10 pages. <https://doi.org/10.1145/3219819.3219885>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

KDD '18, August 19-23, 2018, London, United Kingdom

© 2018 Association for Computing Machinery.

ACM ISBN 978-1-4503-5552-0/18/08...\$15.00

<https://doi.org/10.1145/3219819.3219885>

## 1 INTRODUCTION

During last decade Search architectures, which were typically based on classic Information Retrieval, have seen an increased presence of Machine Learning in its various components [2], especially in Search Ranking which often has challenging objectives depending on the type of content that is being searched over. The main reason behind this trend is the rise in the amount of search data that can be collected and analyzed. The large amounts of collected data open up possibilities for using Machine Learning to personalize search results for a particular user based on previous searches and recommend similar content to recently consumed one.

The objective of any search algorithm can vary depending on the platform at hand. While some platforms aim at increasing website engagement (e.g. clicks and time spent on news articles that are being searched), others aim at maximizing conversions (e.g. purchases of goods or services that are being searched over), and in the case of two sided marketplaces we often need to optimize the search results for both sides of the marketplace, i.e. sellers and buyers. The two sided marketplaces have emerged as a viable business model in many real world applications. In particular, we have moved from the social network paradigm to a network with two distinct types of participants representing supply and demand. Example industries include accommodation (Airbnb), ride sharing (Uber, Lyft), online shops (Etsy), etc. Arguably, content discovery and search ranking for these types of marketplaces need to satisfy both supply and demand sides of the ecosystem in order to grow and prosper.

In the case of Airbnb, there is a clear need to optimize search results for both hosts and guests, meaning that given an input query with location and trip dates we need to rank high listings whose location, price, style, reviews, etc. are appealing to the guest and, at the same time, are a good match in terms of host preferences for trip duration and lead days. Furthermore, we need to detect listings that would likely reject the guest due to bad reviews, pets, length of stay, group size or any other factor, and rank these listings lower. To achieve this we resort to using Learning to Rank. Specifically, we formulate the problem as pairwise regression with positive utilities for bookings and negative utilities for rejections, which we optimize using a modified version of Lambda Rank [4] model that jointly optimizes ranking for both sides of the marketplace.

Since guests typically conduct multiple searches before booking, i.e. click on more than one listing and contact more than one host during their search session, we can use these in-session signals, i.e. clicks, host contacts, etc. for Real-time Personalization where the aim is to show to the guest more of the listings similar to the ones we think they liked since starting the search session. At the same time we can use the negative signal, e.g. skips of high ranked listings, to show to the guest less of the listings similar to the ones we think

they did not like. To be able to calculate similarities between listings that guest interacted with and candidate listings that need to be ranked we propose to use listing embeddings, low-dimensional vector representations learned from search sessions. We leverage these similarities to create personalization features for our Search Ranking Model and to power our Similar Listing Recommendations, the two platforms that drive 99% of bookings at Airbnb.

In addition to Real-time Personalization using immediate user actions, such as clicks, that can be used as proxy signal for short-term user interest, we introduce another type of embeddings trained on bookings to be able to capture user's long-term interest. Due to the nature of travel business, where users travel 1-2 times per year on average, bookings are a sparse signal, with a long tail of users with a single booking. To tackle this we propose to train embeddings at a level of user type, instead of a particular user id, where type is determined using many-to-one rule-based mapping that leverages known user attributes. At the same time we learn listing type embeddings in the same vector space as user type embeddings. This enables us to calculate similarities between user type embedding of the user who is conducting a search and listing type embeddings of candidate listings that need to be ranked.

Compared to previously published work on embeddings for personalization on the Web, novel contributions of this paper are:

- **Real-time Personalization** - Most of the previous work on personalization and item recommendations using embeddings [8, 11] is deployed to production by forming tables of user-item and item-item recommendations offline, and then reading from them at the time of recommendation. We implemented a solution where embeddings of items that user most recently interacted with are combined in an online manner to calculate similarities to items that need to be ranked.
- **Adapting Training for Congregated Search** - Unlike in Web search, the search on travel platforms is often congregated, where users frequently search only within a certain market, e.g. Paris., and rarely across different markets. We adapted the embedding training algorithm to take this into account when doing negative sampling, which lead to capturing better within-market listings similarities.
- **Leveraging Conversions as Global Context** - We recognize the importance of click sessions that end up in conversion, in our case booking. When learning listing embeddings we treat the booked listing as global context that is always being predicted as the window moves over the session.
- **User Type Embeddings** - Previous work on training user embeddings to capture their long-term interest [6, 27] train a separate embedding for each user. When target signal is sparse, there is not enough data to train a good embedding representation for each user. Not to mention that storing embeddings for each user to perform online calculations would require lot of memory. For that reason we propose to train embeddings at a level of user type, where groups of users with same type will have the same embedding.
- **Rejections as Explicit Negatives** - To reduce recommendations that result in rejections we encode host preference signal in user and listing type embeddings by treating host rejections as explicit negatives during training.

For short-term interest personalization we trained listing embeddings using more than 800 million search clicks sessions, resulting in high quality listing representations. We used extensive offline and online evaluation on real search traffic which showed that adding embedding features to the ranking model resulted in significant booking gain. In addition to the search ranking algorithm, listing embeddings were successfully tested and launched for similar listing recommendations where they outperformed the existing algorithm click-through rate (CTR) by 20%.

For long-term interest personalization we trained user type and listing type embeddings using sequences of booked listings by 50 million users. Both user and listing type embeddings were learned in the same vector space, such that we can calculate similarities between user type and listing types of listings that need to be ranked. The similarity was used as an additional feature for search ranking model and was also successfully tested and launched.

## 2 RELATED WORK

In a number of Natural Language Processing (NLP) applications classic methods for language modeling that represent words as high-dimensional, sparse vectors have been replaced by Neural Language models that learn word embeddings, i.e. low-dimensional representations of words, through the use of neural networks [25, 27]. The networks are trained by directly taking into account the word order and their co-occurrence, based on the assumption that words frequently appearing together in the sentences also share more statistical dependence. With the development of highly scalable continuous bag-of-words (CBOW) and skip-gram (SG) language models for word representation learning [17], the embedding models have been shown to obtain state-of-the-art performance on many traditional language tasks after training on large text data.

More recently, the concept of embeddings has been extended beyond word representations to other applications outside of NLP domain. Researchers from the Web Search, E-commerce and Marketplace domains have quickly realized that just like one can train word embeddings by treating a sequence of words in a sentence as context, same can be done for training embeddings of user actions, e.g. items that were clicked or purchased [11, 18], queries and ads that were clicked [8, 9], by treating sequence of user actions as context. Ever since, we have seen embeddings being leveraged for various types of recommendations on the Web, including music recommendations [26], job search [13], app recommendations [21], movie recommendations [3, 7], etc. Furthermore, it has been shown that items which user interacted with can be leveraged to directly learn user embeddings in the same feature space as item embeddings, such that direct user-item recommendations can be made [6, 10, 11, 24, 27]. Alternative approach, specifically useful for cold-start recommendations, is to still to use text embeddings (e.g. ones publicly available at <https://code.google.com/p/word2vec>) and leverage item and/or user meta data (e.g. title and description) to compute their embeddings [5, 14, 19, 28]. Finally, similar extensions of embedding approaches have been proposed for Social Network analysis, where random walks on graphs can be used to learn embeddings of nodes in graph structure [12, 20].

Embedding approaches have had a major impact in both academia and industry circles. Recent industry conference publications and

talks show that they have been successfully deployed in various personalization, recommendation and ranking engines of major Web companies, such as Yahoo [8, 11, 29], Etsy [1], Criteo [18], LinkedIn [15, 23], Tinder [16], Tumblr [10], Instacart [22], Facebook [28].

### 3 METHODOLOGY

In the following we introduce the proposed methodology for the task of listing recommendations and listing ranking in search at Airbnb. We describe two distinct approaches, i.e. listing embeddings for short-term real-time personalization and user-type & listing type embeddings for long term personalization, respectively.

#### 3.1 Listing Embeddings

Let us assume we are given a set  $\mathcal{S}$  of  $S$  click sessions obtained from  $N$  users, where each session  $s = (l_1, \dots, l_M) \in \mathcal{S}$  is defined as an uninterrupted sequence of  $M$  listing ids that were clicked by the user. A new session is started whenever there is a time gap of more than 30 minutes between two consecutive user clicks. Given this data set, the aim is to learn a  $d$ -dimensional real-valued representation  $\mathbf{v}_{l_i} \in \mathbb{R}^d$  of each unique listing  $l_i$ , such that similar listings lie nearby in the embedding space.

More formally, the objective of the model is to learn listing representations using the skip-gram model [17] by maximizing the objective function  $\mathcal{L}$  over the entire set  $\mathcal{S}$  of search sessions, defined as follows

$$\mathcal{L} = \sum_{s \in \mathcal{S}} \sum_{l_i \in s} \left( \sum_{-m \leq j \leq m, i \neq j} \log \mathbb{P}(l_{i+j} | l_i) \right), \quad (1)$$

Probability  $\mathbb{P}(l_{i+j} | l_i)$  of observing a listing  $l_{i+j}$  from the contextual neighborhood of clicked listing  $l_i$  is defined using the soft-max

$$\mathbb{P}(l_{i+j} | l_i) = \frac{\exp(\mathbf{v}_{l_i}^\top \mathbf{v}'_{l_{i+j}})}{\sum_{l' \in \mathcal{V}} \exp(\mathbf{v}_{l_i}^\top \mathbf{v}'_{l'})}, \quad (2)$$

where  $\mathbf{v}_l$  and  $\mathbf{v}'_l$  are the input and output vector representations of listing  $l$ , hyperparameter  $m$  is defined as a length of the relevant forward looking and backward looking context (neighborhood) for a clicked listing, and  $\mathcal{V}$  is a vocabulary defined as a set of unique listings ids in the data set. From (1) and (2) we see that the proposed approach models temporal context of listing click sequences, where listings with similar contexts (i.e., with similar neighboring listings in search sessions) will have similar representations.

Time required to compute gradient  $\nabla \mathcal{L}$  of the objective function in (1) is proportional to the vocabulary size  $|\mathcal{V}|$ , which for large vocabularies, e.g. several millions listing ids, is an infeasible task. As an alternative we used negative sampling approach proposed in [17], which significantly reduces computational complexity. Negative sampling can be formulated as follows. We generate a set  $\mathcal{D}_p$  of *positive pairs*  $(l, c)$  of clicked listings  $l$  and their contexts  $c$  (i.e., clicks on other listings by the same user that happened before and after click on listing  $l$  within a window of length  $m$ ), and a set  $\mathcal{D}_n$  of *negative pairs*  $(l, c)$  of clicked listings and  $n$  randomly sampled listings from the entire vocabulary  $\mathcal{V}$ . The optimization objective then becomes

$$\arg \max_{\theta} \sum_{(l, c) \in \mathcal{D}_p} \log \frac{1}{1 + e^{-\mathbf{v}'_c \mathbf{v}_l}} + \sum_{(l, c) \in \mathcal{D}_n} \log \frac{1}{1 + e^{\mathbf{v}'_c \mathbf{v}_l}}, \quad (3)$$




Figure 1: Skip-gram model for Listing Embeddings

where parameters  $\theta$  to be learned are  $\mathbf{v}_l$  and  $\mathbf{v}_c$ ,  $l, c \in \mathcal{V}$ . The optimization is done via stochastic gradient ascent.

**Booked Listing as Global Context.** We can break down the click sessions set  $\mathcal{S}$  into 1) *booked sessions*, i.e. click sessions that end with user booking a listing to stay at, and 2) *exploratory sessions*, i.e. click sessions that do not end with booking, i.e. users were just browsing. Both are useful from the standpoint of capturing contextual similarity, however *booked sessions* can be used to adapt the optimization such that at each step we predict not only the neighboring clicked listings but the eventually booked listing as well. This adaptation can be achieved by adding booked listing as *global context*, such that it will always be predicted no matter if it is within the context window or not. Consequently, for *booked sessions* the embedding update rule becomes

$$\arg \max_{\theta} \sum_{(l, c) \in \mathcal{D}_p} \log \frac{1}{1 + e^{-\mathbf{v}'_c \mathbf{v}_l}} + \sum_{(l, c) \in \mathcal{D}_n} \log \frac{1}{1 + e^{\mathbf{v}'_c \mathbf{v}_l}} + \log \frac{1}{1 + e^{-\mathbf{v}'_{l_b} \mathbf{v}_l}}, \quad (4)$$

where  $\mathbf{v}_{l_b}$  is the embedding of the booked listing  $l_b$ . For *exploratory sessions* the updates are still conducted by optimizing objective (3).

Figure 1 shows a graphical representation of how listing embeddings are learned from *booked sessions* using a sliding window of size  $2n + 1$  that slides from the first clicked listing to the booked listing. At each step the embedding of the central listing  $\mathbf{v}_l$  is being updated such that it predicts the embeddings of the context listings  $\mathbf{v}_c$  from  $\mathcal{D}_p$  and the booked listing  $\mathbf{v}_{l_b}$ . As the window slides some listings fall in and out of the context set, while the booked listing always remains within it as global context (dotted line).

**Adapting Training for Congregated Search.** Users of online travel booking sites typically search only within a single market, i.e. location they want to stay at. As a consequence, there is a high probability that  $\mathcal{D}_p$  contains listings from the same market. On the other hand, due to random sampling of negatives, it is very likely that  $\mathcal{D}_n$  contains mostly listings that are not from the same markets as listings in  $\mathcal{D}_p$ . At each step, for a given central listing  $l$ , the positive context mostly consist of listings from the same market as  $l$ , while the negative context mostly consists of listings that are not from the same market as  $l$ . We found that this imbalance leads to learning sub-optimal within-market similarities. To address this issue we propose to add a set of random negatives  $\mathcal{D}_{mn}$ , sampled from the market of the central listing  $l$ ,

$$\arg \max_{\theta} \sum_{(l, c) \in \mathcal{D}_p} \log \frac{1}{1 + e^{-\mathbf{v}'_c \mathbf{v}_l}} + \sum_{(l, c) \in \mathcal{D}_n} \log \frac{1}{1 + e^{\mathbf{v}'_c \mathbf{v}_l}} + \log \frac{1}{1 + e^{-\mathbf{v}'_{l_b} \mathbf{v}_l}} + \sum_{(l, m_n) \in \mathcal{D}_{mn}} \log \frac{1}{1 + e^{\mathbf{v}'_{m_n} \mathbf{v}_l}}. \quad (5)$$

where parameters  $\theta$  to be learned are  $\mathbf{v}_l$  and  $\mathbf{v}_c$ ,  $l, c \in \mathcal{V}$ .




Figure 2: California Listing Embedding Clusters



Figure 3: Similar Listings using Embeddings

**Cold start listing embeddings.** Every day new listings are created by hosts and made available on Airbnb. At that point these listings do not have an embedding because they were not present in the click sessions  $\mathcal{S}$  training data. To create embeddings for new listings we propose to utilize existing embeddings of other listings.

Upon listing creation the host is required to provide information about the listing, such as location, price, listing type, etc. We use the provided meta-data about the listing to find 3 geographically closest listings (within a 10 miles radius) that have embeddings, are of same listing type as the new listing (e.g. Private Room) and belong to the same price bucket as the new listing (e.g. \$20 – \$25 per night). Next, we calculate the mean vector using 3 embeddings of identified listings to form the new listing embedding. Using this technique we are able to cover more than 98% of new listings.

**Examining Listing Embeddings.** To evaluate what characteristics of listings were captured by embeddings we examine the  $d = 32$  dimensional embeddings trained using (5) on 800 million click sessions. First, by performing  $k$ -means clustering on learned embeddings we evaluate if geographical similarity is encoded. Figure 2, which shows resulting 100 clusters in California, confirms that listings from similar locations are clustered together. We found the clusters very useful for re-evaluating our definitions of travel markets. Next, we evaluate average cosine similarities between

### Embedding Evaluation Tool

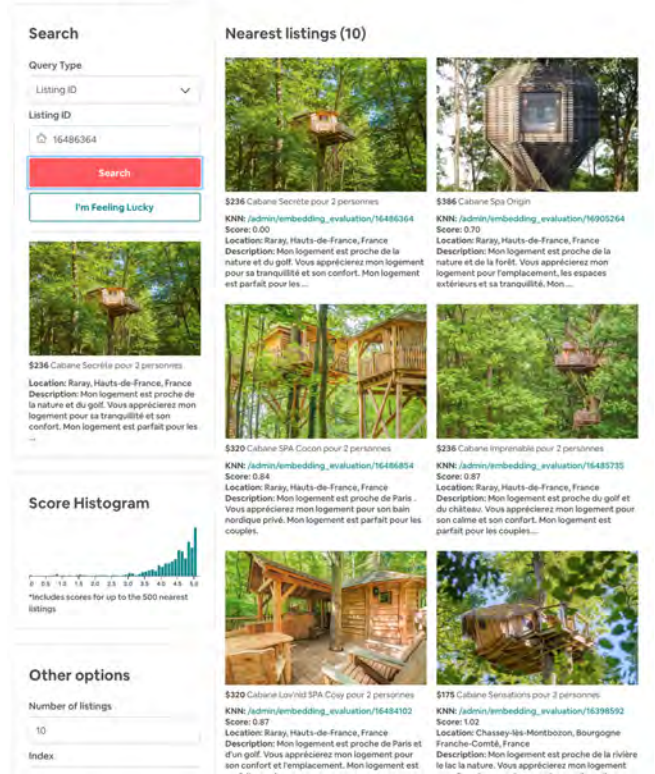


Figure 4: Embeddings Evaluation Tool

listings from Los Angeles of different listing types (Table 1) and between listings of different price ranges (Table 2). From those tables it can be observed that cosine similarities between listings of same type and price ranges are much higher compared to similarities between listings of different types and price ranges. Therefore, we can conclude that those two listing characteristics are well encoded in the learned embeddings as well.

While some listing characteristics, such as price, do not need to be learned because they can be extracted from listing meta-data, other types of listing characteristics, such as architecture, style and feel are much harder to extract in form of listing features. To evaluate if these characteristics are captured by embeddings we can examine  $k$ -nearest neighbors of unique architecture listings in the listing embedding space. Figure 3 shows one such case where for a listing of unique architecture on the left, the most similar listings are of the same style and architecture. To be able to conduct fast and easy explorations in the listing embedding space we developed an internal **Similarity Exploration Tool** shown in Figure 4.

**Demonstration video** of this tool, which is available online at <https://youtu.be/1kJSAG91TrI>, shows many more examples of embeddings being able to find similar listings of the same unique architecture, including houseboats, treehouses, castles, chalets, beach-front apartments, etc.

**Table 1: Cosine similarities between different Listing Types**

Room Type	Entire Home	Private Room	Shared Room
Entire Home	<b>0.895</b>	0.875	0.848
Private Room		<b>0.901</b>	0.865
Shared Room			<b>0.896</b>

**Table 2: Cosine similarities between different Price Ranges**

Price Range	<\$30	\$30-\$60	\$60-\$90	\$90-\$120	\$120+
<\$30	<b>0.916</b>	0.887	0.882	0.871	0.854
\$30-\$60		<b>0.906</b>	0.889	0.876	0.865
\$60-\$90			<b>0.902</b>	0.883	0.880
\$90-\$120				<b>0.898</b>	0.890
\$120+					<b>0.909</b>

### 3.2 User-type & Listing-type Embeddings

Listing embeddings described in Section 3.1. that were trained using click sessions are very good at finding similarities between listings of the same market. As such, they are suitable for short-term, in-session, personalization where the aim is to show to the user listings that are similar to the ones they clicked during the immanent search session.

However, in addition to in-session personalization, based on signals that just happened within the same session, it would be useful to personalize search based on signals from user’s longer-term history. For example, given a user who is currently searching for a listing in Los Angeles, and has made past bookings in New York and London, it would be useful to recommend listings that are similar to those previously booked ones.

While some cross-market similarities are captured in listing embeddings trained using clicks, a more principal way of learning such cross-market similarities would be to learn from sessions constructed of listings that a particular user booked over time. Specifically, let us assume we are given a set  $\mathcal{S}_b$  of booking sessions obtained from  $N$  users, where each booking session  $s_b = (l_{b1}, \dots, l_{bM})$  is defined as a sequence of listings booked by user  $j$  ordered in time. Attempting to learn embeddings  $\mathbf{v}_{l_{id}}$  for each  $listing\_id$  using this type of data would be challenging in many ways:

- First, booking sessions data  $\mathcal{S}_b$  is much smaller than click sessions data  $\mathcal{S}$  because bookings are less frequent events.
- Second, many users booked only a single listing in the past and we cannot learn from a session of length 1.
- Third, to learn a meaningful embedding for any entity from contextual information at least 5 – 10 occurrences of that entity are needed in the data, and there are many  $listing\_ids$  on the platform that were booked less than 5 – 10 times.
- Finally, long time intervals may pass between two consecutive bookings by the user, and in that time user preferences, such as price point, may change, e.g. due to career change.

To address these very common marketplace problems in practice, we propose to learn embeddings at a level of  $listing\_type$  instead of  $listing\_id$ . Given meta-data available for a certain  $listing\_id$  such as location, price, listing type, capacity, number of beds, etc., we use a

**Table 3: Mappings of listing meta data to listing type buckets**

Buckets	1	2	3	4	5	6	7	8
Country	US	CA	GB	FR	MX	AU	ES	...
Listing Type	Ent	Priv	Share					
\$ per Night	<40	40-55	56-69	70-83	84-100	101-129	130-189	190+
\$ per Guest	<21	21-27	28-34	35-42	43-52	53-75	76+	
Num Reviews	0	1	2-5	6-10	11-35	35+		
Listing 5 Star %	0-40	41-60	61-90	90+				
Capacity	1	2	3	4	5	6+		
Num Beds	1	2	3	4+				
Num Bedrooms	0	1	2	3	4+			
Num Bathroom	0	1	2	3+				
New Guest Acc %	<60	61-90	>91					

**Table 4: Mappings of user meta data to user type buckets**

Buckets	1	2	3	4	5	6	7	8
Market	SF	NYC	LA	HK	PHL	AUS	LV	...
Language	en	es	fr	jp	ru	ko	de	...
Device Type	Mac	Msft	Andr	Ipad	Tablet	Iphone		
Full Profile	Yes	No						
Profile Photo	Yes	No						
Num Bookings	0	1	2-7	8+				
\$ per Night	<40	40-55	56-69	70-83	84-100	101-129	130-189	190+
\$ per Guest	<21	21-27	28-34	35-42	43-52	53-75	76+	
Capacity	<2	2-2.6	2.7-3	3.1-4	4.1-6	6.1+		
Num Reviews	<1	1-3.5	3.6-10	>10				
Listing 5 Star %	0-40	41-60	61-90	90+				
Guest 5 Star %	0-40	41-60	61-90	90+				

rule-based mapping defined in Table 3 to determine its  $listing\_type$ . For example, an *Entire Home* listing from *US* that has a 2 person capacity, 1 bed, 1 bedroom & 1 bathroom, with Average Price Per Night of \$60.8, Average Price Per Night Per Guest of \$29.3, 5 reviews, all 5 stars, and 100% *New Guest Accept Rate* would map into  $listing\_type = US\_lt1\_pn3\_pg3\_r3\_5s4\_c2\_b1\_bd2\_bt2\_nu3$ . Buckets are determined in a data-driven manner to maximize for coverage in each  $listing\_type$  bucket. The mapping from  $listing\_id$  to a  $listing\_type$  is a many-to-one mapping, meaning that many listings will map into the same  $listing\_type$ .

To account for user ever-changing preferences over time we propose to learn  $user\_type$  embeddings in the same vector space as  $listing\_type$  embeddings. The  $user\_type$  is determined using a similar procedure we applied to listings, i.e. by leveraging meta-data about user and their previous bookings, defined in Table 4. For example, for a user from San Francisco with MacBook laptop, English language settings, full profile with user photo, 83.4% average Guest 5 star rating from hosts, who has made 3 bookings in the past, where the *average* statistics of booked listings were \$52.52 Price Per Night, \$31.85 Price Per Night Per Guest, 2.33 Capacity, 8.24 Reviews and 76.1% Listing 5 star rating, the resulting  $user\_type$  is  $SF\_lg1\_dt1\_fp1\_pp1\_nb1\_ppn2\_ppg3\_c2\_nr3\_l5s3\_g5s3$ . When generating booking sessions for training embeddings we calculate the  $user\_type$  up to the latest booking. For users who made their first booking  $user\_type$  is calculated based on the first 5 rows from Table 4 because at the time of booking we had no prior information about past bookings. This is convenient, because learned embeddings for  $user\_types$  which are based on first 5 rows can be used for cold-start personalization for logged-out users and new users with no past bookings.

**Training Procedure.** To learn  $user\_type$  and  $listing\_type$  embeddings in the same vector space we incorporate the  $user\_type$



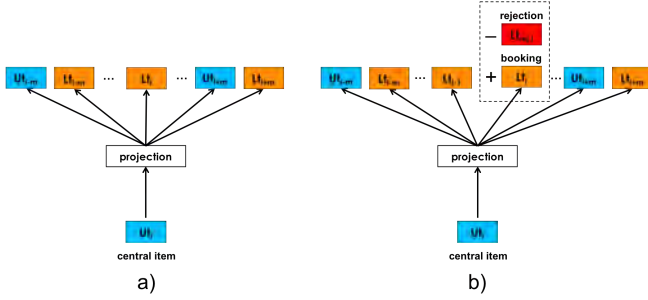


Figure 5: Listing Type and User Type Skip-gram model

into the booking sessions. Specifically, we form a set  $\mathcal{S}_b$  consisting of  $N_b$  booking sessions from  $N$  users, where each session  $s_b = (u_{type_1} l_{type_1}, \dots, u_{type_M} l_{type_M}) \in \mathcal{S}_b$  is defined as a sequence of booking events, i.e.  $(user\_type, listing\_type)$  tuples ordered in time. Note that each session consists of bookings by same  $user\_id$ , however for a single  $user\_id$  their  $user\_types$  can change over time, similarly to how  $listing\_types$  for the same listing can change over time as they receive more bookings.

The objective that needs to be optimized can be defined similarly to (3), where instead of listing  $l$ , the center item that needs to be updated is either  $user\_type$  ( $u_t$ ) or  $listing\_type$  ( $l_t$ ) depending on which one is caught in the sliding window. For example, to update the central item which is a  $user\_type$  ( $u_t$ ) we use

$$\arg\max_{\theta} \sum_{(u_t, c) \in \mathcal{D}_{book}} \log \frac{1}{1 + e^{-v_c^T v_{u_t}}} + \sum_{(u_t, c) \in \mathcal{D}_{neg}} \log \frac{1}{1 + e^{v_c^T v_{u_t}}}, \quad (6)$$

where  $\mathcal{D}_{book}$  contains the  $user\_type$  and  $listing\_type$  from recent user history, specifically user bookings from near past and near future with respect to central item's timestamp, while  $\mathcal{D}_{neg}$  contains random  $user\_type$  or  $listing\_type$  instances used as negatives. Similarly, if the central item is a  $listing\_type$  ( $l_t$ ) we optimize the following objective

$$\arg\max_{\theta} \sum_{(l_t, c) \in \mathcal{D}_{book}} \log \frac{1}{1 + e^{-v_c^T v_{l_t}}} + \sum_{(l_t, c) \in \mathcal{D}_{neg}} \log \frac{1}{1 + e^{v_c^T v_{l_t}}}. \quad (7)$$

Figure 5a (on the left) shows a graphical representation of this model, where central item represents  $user\_type$  ( $u_t$ ) for which the updates are performed as in (6).

Since *booking sessions* by definition mostly contain listings from different markets, there is no need to sample additional negatives from same market as the booked listing, like we did in Session 3.1. to account for the congregated search in *click sessions*.

**Explicit Negatives for Rejections.** Unlike clicks that only reflect guest-side preferences, bookings reflect host-side preferences as well, as there exists an explicit feedback from the host, in form of accepting guest's request to book or rejecting guest's request to book. Some of the reasons for host rejections are bad guest star ratings, incomplete or empty guest profile, no profile picture, etc. These characteristics are part of  $user\_type$  definition from Table 4.

Host rejections can be utilized during training to encode the host preference signal in the embedding space in addition to the guest preference signal. The whole purpose of incorporating the rejection signal is that some  $listing\_types$  are less sensitive to  $user\_types$  with no bookings, incomplete profiles and less than average guest

Table 5: Recommendations based on type embeddings

User Type	
$SF\_lg1\_dt1\_fp1\_pp1\_nb3\_ppn5\_ppg5\_c4\_nr3\_l5s3\_g5s3$	
Listing Type	Sim
$US\_lt1\_pn4\_pg5\_r5\_5s4\_c2\_b1\_bd3\_bt3\_nu3$ (large, good reviews)	0.629
$US\_lt1\_pn3\_pg3\_r5\_5s2\_c3\_b1\_bd2\_bt2\_nu3$ (cheaper, bad reviews)	0.350
$US\_lt2\_pn3\_pg3\_r5\_5s4\_c1\_b1\_bd2\_bt2\_nu3$ (priv room, good reviews)	0.241
$US\_lt2\_pn2\_pg2\_r5\_5s2\_c1\_b1\_bd2\_bt2\_nu3$ (cheaper, bad reviews)	0.169
$US\_lt3\_pn1\_pg1\_r5\_5s3\_c1\_b1\_bd2\_bt2\_nu3$ (shared room, bad reviews)	0.121

star ratings than others, and we want the embeddings of those  $listing\_types$  and  $user\_types$  to be closer in the vector space, such that recommendations based on embedding similarities would reduce future rejections in addition to maximizing booking chances.

We formulate the use of the rejections as explicit negatives in the following manner. In addition to sets  $\mathcal{D}_{book}$  and  $\mathcal{D}_{neg}$ , we generate a set  $\mathcal{D}_{rej}$  of pairs of  $(u_t, l_t)$  of  $user\_type$  or  $listing\_type$  that were involved in a rejection event. As depicted in Figure 5b (on the right), we specifically focus on the cases when host rejections (labeled with a minus sign) were followed by a successful booking (labeled with a plus sign) of another listing by the same user. The new optimization objective can then be formulated as

$$\arg\max_{\theta} \sum_{(u_t, c) \in \mathcal{D}_{book}} \log \frac{1}{1 + \exp^{-v_c^T v_{u_t}}} + \sum_{(u_t, c) \in \mathcal{D}_{neg}} \log \frac{1}{1 + \exp^{v_c^T v_{u_t}}} + \sum_{(u_t, l_t) \in \mathcal{D}_{reject}} \log \frac{1}{1 + \exp^{v_{l_t}^T v_{u_t}}}. \quad (8)$$

in case of updating the central item which is a  $user\_type$  ( $u_t$ ), and

$$\arg\max_{\theta} \sum_{(l_t, c) \in \mathcal{D}_{book}} \log \frac{1}{1 + \exp^{-v_c^T v_{l_t}}} + \sum_{(l_t, c) \in \mathcal{D}_{neg}} \log \frac{1}{1 + \exp^{v_c^T v_{l_t}}} + \sum_{(l_t, u_t) \in \mathcal{D}_{reject}} \log \frac{1}{1 + \exp^{v_{u_t}^T v_{l_t}}}. \quad (9)$$

in case of updating the central item which is a  $listing\_type$  ( $l_t$ ).

Given learned embeddings for all  $user\_types$  and  $listing\_types$ , we can recommend to the user the most relevant listings based on the cosine similarities between user's current  $user\_type$  embedding and  $listing\_type$  embeddings of candidate listings. For example, in Table 5 we show cosine similarities between  $user\_type = SF\_lg1\_dt1\_fp1\_pp1\_nb3\_ppn5\_ppg5\_c4\_nr3\_l5s3\_g5s3$  who typically books high quality, spacious listings with lots of good reviews and several different  $listing\_types$  in US. It can be observed that listing types that best match these user preferences, i.e. entire home, lots of good reviews, large and above average price, have high cosine similarity, while the ones that do not match user preferences, i.e. ones with less space, lower price and small number of reviews have low cosine similarity.